

Title: Data7202 A2 Report

Name: Xinqian Wang


ID: s4565489

Tips: the code screenshot will be attached in the Appendix. For code script will also be attached in the folder.

● 1.

Question: Apply Principal Component Regression (PCR) with all possible number of principal components. Using the 10-Fold Cross-Validation, plot the mean squared error as a function of the number of components and determine the optimal number of components.


Answer:



The optimal number of components is 18.

Question: Apply the Lasso method and plot the 10-Fold Cross-Validation mean squared error as a function of λ . Determine the best λ and the corresponding mean squared error.

Answer:



The best λ should be the alpha equals to 5.

● 2.

Question: Construct a Poisson regression model and report the coefficients (for type, construction, operation, and months), and the corresponding 95% CIs.

Answer:

By using the code `model = sm.GLM(Y,X,family=sm.families.Poisson())`, we build a Poisson Regression Model. The coefficients for **type**, **construction**, **operation**, and **months** are below:

Generalized Linear Model Regression Results						
=====						
Dep. Variable:	damage	No. Observations:	34			
Model:	GLM	Df Residuals:	30			
Model Family:	Poisson	Df Model:	3			
Link Function:	log	Scale:	1.0000			
Method:	IRLS	Log-Likelihood:	-145.96			
Date:	Mon, 12 Apr 2021	Deviance:	194.06			
Time:	13:30:54	Pearson chi2:	178.			
No. Iterations:	6					
Covariance Type:	nonrobust					
=====						
	coef	std err	z	P> z	[0.025	0.975]

type	-0.2237	0.048	-4.693	0.000	-0.317	-0.130
construction	0.3714	0.060	6.231	0.000	0.255	0.488
operation	0.7680	0.103	7.471	0.000	0.567	0.969
months	8.095e-05	2.84e-06	28.487	0.000	7.54e-05	8.65e-05
=====						

● 3.

Question: Compute the multiple regression of Time on Cases and Distance. State the fitted model, the estimated residual standard deviation, and the P-values for the overall model and each of the two predictors.

Answer:

```
Call:
lm(formula = Time ~ Cases + Distance, data = data)

Residuals:
    Min       1Q   Median       3Q      Max
-5.7880 -0.6629  0.4364  1.1566  7.4197

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  2.341231   1.096730   2.135 0.044170 *
Cases        1.615907   0.170735   9.464 3.25e-09 ***
Distance     0.014385   0.003613   3.981 0.000631 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.259 on 22 degrees of freedom
Multiple R-squared:  0.9596,    Adjusted R-squared:  0.9559
F-statistic: 261.2 on 2 and 22 DF,  p-value: 4.687e-16
```

The Residual standard error $= \sqrt{\frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{df}} = 3.259$ on 22 degrees of freedom.

The P-value of overall model is 4.687e-16.


The P-value of the attribute Cases is 3.25e-09.

The P-value of the attribute Distance is 0.000631.

Question: Obtain residual plots and the histogram of the residuals. Comment on these.


Answer:

Residual plot is as below:




Residual histogram is as below:


Histogram of residuals(fit)




Residuals vs Fitted




Normal Q-Q



Scale-Location



Residuals vs Leverage



The **histogram** shows most of the residuals have the values range between -2 and 2.

The **Residuals vs Fitted plot** shows whether the residuals have non-linear patterns. The plot fails to show equally spread residuals around a horizontal line without distinct patterns, which indicates that the non-linear relationship was not explained by the model and was left out in the residuals.

The **QQ plot** is used to show if the residuals are normally distributed or not. The plot just tells that the residuals are probably not following a normal distribution.

The **Spread-Location plot** shows if residuals are spread equally along the ranges of predictors. The plot shows that most of the residuals are spread in the left side of the plot whereas only a few points distributed on the right. Especially the 9th point, it's value seems not common as other points.

This **Residuals vs Leverage plot** helps us to find influential cases if any. Not all outliers are influential in linear regression analysis (whatever outliers mean). It is usually the case that a point is judged as an outlier if its cook's distance is greater than 0.5. In this case, the point 22 and point 9 would interest us to do further research, which also means they are more influential than any other points when the regression model is decided.

Question: There is an observation in this data set which is extremely influential according to Cook's distance. Which observation is it? Display a Cook's distance plot to determine the Cook's distance of the next most influential observation.

Answer:

The 9-th point is the most influential point, according to the Influence table below and the plots above. The Cook's distance plot is placed on the former plot. The next most influential observation would be the 22-ed point.

```

Influence measures of
lm(formula = Time ~ Cases + Distance, data = data) :

      dfb.1_ dfb.Cass dfb.Dstn  dffit cov.r  cook.d  hat inf
1 -0.18727  0.41131 -0.43486 -0.5709 0.871 1.00e-01 0.1018
2  0.08979 -0.04776  0.01441  0.0986 1.215 3.38e-03 0.0707
3 -0.00352  0.00395 -0.00285 -0.0052 1.276 9.46e-06 0.0987
4  0.45196  0.08828 -0.27337  0.5008 0.876 7.76e-02 0.0854
5 -0.03167 -0.01330  0.02424 -0.0395 1.240 5.43e-04 0.0750
6 -0.01468  0.00179  0.00108 -0.0188 1.200 1.23e-04 0.0429
7  0.07807 -0.02228 -0.01102  0.0790 1.240 2.17e-03 0.0818
8  0.07120  0.03338 -0.05382  0.0938 1.206 3.05e-03 0.0637
9 -2.57574  0.92874  1.50755  4.2961 0.342 3.42e+00 0.4983 *
10 0.10792 -0.33816  0.34133  0.3987 1.305 5.38e-02 0.1963
11 -0.03427  0.09253 -0.00269  0.2180 1.172 1.62e-02 0.0861
12 -0.03027 -0.04867  0.05397 -0.0677 1.291 1.60e-03 0.1137
13  0.07237 -0.03562  0.01134  0.0813 1.207 2.29e-03 0.0611
14  0.04952 -0.06709  0.06182  0.0974 1.228 3.29e-03 0.0782
15  0.02228 -0.00479  0.00684  0.0426 1.192 6.32e-04 0.0411
16 -0.00269  0.06442 -0.08419 -0.0972 1.369 3.29e-03 0.1659
17  0.02886  0.00649 -0.01570  0.0339 1.219 4.01e-04 0.0594
18  0.24856  0.18973 -0.27243  0.3653 1.069 4.40e-02 0.0963
19  0.17256  0.02357 -0.09897  0.1862 1.215 1.19e-02 0.0964
20  0.16804 -0.21500 -0.09292 -0.6718 0.760 1.32e-01 0.1017
21 -0.16193 -0.29718  0.33641 -0.3885 1.238 5.09e-02 0.1653
22  0.39857 -1.02541  0.57314 -1.1950 1.398 4.51e-01 0.3916 *
23 -0.15985  0.03729 -0.05265 -0.3075 0.890 2.99e-02 0.0413
24 -0.11972  0.40462 -0.46545 -0.5711 0.948 1.02e-01 0.1206
25 -0.01682  0.00085  0.00559 -0.0176 1.231 1.08e-04 0.0666

```

● 4.

Question: Derive the posterior distribution of θ .

Answer:

Firstly, calculating the likelihood:

$$p(y | \theta) = \frac{1}{\theta}$$

Secondly, calculating the prior:

$$\mathbf{p}(\theta | \alpha, x_m) = \begin{cases} \frac{\alpha x_m^\alpha}{\theta^{\alpha+1}} & \theta \geq x_m \\ 0 & \theta < x_m \end{cases}$$

Thirdly, calculating the posterior and prove it proportional to a Pareto distribution:

$$\begin{aligned}
P(\theta \mid Y) &= \frac{P(Y \mid \theta) \cdot P(\theta)}{P(Y)} \\
&\propto P(Y \mid \theta) \cdot P(\theta) \\
&= \left(\prod_{i=1}^n \mathbf{p}(y_i \mid \theta) \right) \cdot \mathbf{p}(\theta \mid \alpha, x_m) \\
&= \frac{1}{\theta^n} \cdot \mathbf{p}(\theta \mid \alpha, x_m) \\
&= \begin{cases} \frac{\alpha x_m^\alpha}{\theta^{(n+\alpha)+1}} & \theta \geq x_m \\ 0 & \theta < x_m \end{cases} \\
&= \frac{\alpha x_m^{-n}}{n + \alpha} \cdot \begin{cases} \frac{(n+\alpha)x_m^{(n+\alpha)}}{\theta^{(n+\alpha)+1}} & \theta \geq x_m \\ 0 & \theta < x_m \end{cases} \\
&\propto \mathbf{Pareto}(\theta \mid (n + \alpha), x_m)
\end{aligned}$$

● 5.

Question: Find the conditional pdf of X given $Y = y$, and the conditional pdf of Y given $X = x$.

Answer:

For any y such that $f_Y(y) > 0$, the conditional pdf of X given that $Y = y$ is the function of x denoted by $f(x | y)$ and defined by

$$f(y | x) = \frac{f(x,y)}{f_X(x)}$$

$$f(x | y) = \frac{f(x,y)}{f_Y(y)}$$

$$\forall x \geq 0, y \geq 0$$

$$\begin{aligned} f_X(x) &= \int_0^{+\infty} f(x,y) dy \\ &= \int_0^{+\infty} c e^{-(xy+x+y)} dy \\ &= \int_0^{+\infty} c \cdot e^{-x} \cdot e^{-(x+1)y} dy \\ &= c \cdot e^{-x} \cdot \frac{1}{-(1+x)} \cdot \lim_{y \rightarrow +\infty} (e^{-(1+x)y} - 1) \\ &= \frac{c \cdot e^{-x}}{1+x} \end{aligned}$$

Similarly, marginal of Y :

$$f_Y(y) = \frac{c \cdot e^{-y}}{1+y}$$

$$f(y | x) = \frac{f(x,y)}{f_X(x)} = (1+x)e^{-y(1+x)} = \text{Exp}(y; \lambda = (1+x))$$


$$f(x | y) = \frac{f(x,y)}{f_Y(y)} = (1+y)e^{-x(1+y)} = \text{Exp}(x; \lambda = (1+y))$$

Conclusion, for given $Y=y$: $X \sim \text{Exp}(1+y)$
for given $X=x$: $Y \sim \text{Exp}(1+x)$

Question: Write working Python code that implements the Gibbs sampler and outputs 1000 points that are approximately distributed according to f .


Answer:

Attached in the code. The scatter plot as below:



I can also calculate the approximation of the constant c since it is hard to be calculated from the equation of $\int_0^{+\infty} \int_0^{+\infty} f(x,y) dy dx = 1$. By calculating the sum

of the 1000 samples and divided by 1000, we can get a good approximation of c . Repeating it with 1000 times, we can get the range of value c^{-1} . I also calculated the mean value of c^{-1} , which is around **0.069578**. I draw a scatter plot with the value of 1000 number of c^{-1} below:



Appendix

● 1.

```
In [1]: # Package implementation
import pandas as pd
from sklearn.decomposition import PCA
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import KFold
from sklearn.metrics import mean_squared_error
from sklearn.linear_model import LinearRegression
import numpy as np
import matplotlib.pyplot as plt
from sklearn import linear_model

In [2]: df = pd.read_csv("Hitters.csv")
# Change Data Type
labelencoder = LabelEncoder()
df['League'] = labelencoder.fit_transform(df['League'])
df['Division'] = labelencoder.fit_transform(df['Division'])
df['NewLeague'] = labelencoder.fit_transform(df['NewLeague'])

#Generate X, Y
X_cols = [i for i in df.columns if i not in ['Salary']]
Y_cols = [i for i in df.columns if i in ['Salary']]
X = df[X_cols]
Y = df[Y_cols]
```

Q1_a

```
In [3]: def PCR(X, Y, N_Cols, model):
        #PCA
        X = PCA(n_components=N_Cols).fit_transform(X)

        X = pd.DataFrame(X)

        kf = KFold(n_splits=10)
        kf.get_n_splits(X)
        MSE = []

        for train_index, test_index in kf.split(X):
            #print("TRAIN:", train_index, "TEST:", test_index)
            X_train, X_test = X.iloc[train_index, :], X.iloc[test_index, :]
            y_train, y_test = Y.iloc[train_index], Y.iloc[test_index]
            # fit the model
            model.fit(X_train, y_train)
            # predict
            y_pred = model.predict(X_test)
            loss = mean_squared_error(y_test, y_pred)
            MSE.append(loss)
        return np.mean(MSE)

In [4]: model = LinearRegression()
        PCR(X, Y, 19, model)
        Error_List = []
        Component_List = [i for i in range(1, 20)]
        for i in range(1, 20):
            Error_List.append(PCR(X, Y, i, model))


In [5]: plt.xlabel("the total number of components")
        plt.ylabel("the value of MSE")
        plt.scatter([i for i in range(1, 20)], Error_List)
```

Q1_b

```
In [8]: lasso = linear_model.Lasso(fit_intercept=True, alpha=5)
        print(PCR(X, Y, 18, lasso))
        Error_List = []
        Component_List = [i for i in range(1, 20)]
        for i in range(1, 30):
            Error_List.append(PCR(X, Y, 18, linear_model.Lasso(fit_intercept=True, alpha=i)))
        plt.xlabel("the value of alpha")
        plt.ylabel("the value of MSE")
        plt.scatter([i for i in range(1, 30)], Error_List, s = 1)
```

115462.2995924692

Out[8]: <matplotlib.collections.PathCollection at 0x2c64d085e88>



- 2.

```

In [1]: # Package implementation
import pandas as pd
import numpy as np
#from sklearn.preprocessing import OneHotEncoder
#from sklearn.preprocessing import LabelEncoder

import statsmodels.api as sm
from statsmodels.formula.api import ols
from statsmodels.formula.api import poisson
import matplotlib.pyplot as plt

In [2]: df = pd.read_csv("ships.csv")
'''
enc = OneHotEncoder(handle_unknown='ignore')
labelencoder = LabelEncoder()

enc_df_League = pd.DataFrame(enc.fit_transform(df[['type']]).toarray())
df['Type_A'] = labelencoder.fit_transform(enc_df_League[0])
df['Type_B'] = labelencoder.fit_transform(enc_df_League[1])
df['Type_C'] = labelencoder.fit_transform(enc_df_League[2])
df['Type_D'] = labelencoder.fit_transform(enc_df_League[3])
df['Type_E'] = labelencoder.fit_transform(enc_df_League[4])

enc_df_League = pd.DataFrame(enc.fit_transform(df[['construction']]).toarray())
df['Con_1'] = labelencoder.fit_transform(enc_df_League[0])
df['Con_2'] = labelencoder.fit_transform(enc_df_League[1])
df['Con_3'] = labelencoder.fit_transform(enc_df_League[2])
df['Con_4'] = labelencoder.fit_transform(enc_df_League[3])
'''

#df = df.drop(['type', 'construction'], axis=1)

```

```

In [4]: Y = df["damage"]
X = df.drop(['damage'], axis=1)

```

```

In [5]: model = sm.GLM(Y,X,family=sm.families.Poisson())
results = model.fit()
print(results.summary())

```

```

Generalized Linear Model Regression Results
=====
Dep. Variable:          damage    No. Observations:          34
Model:                  GLM      Df Residuals:              30
Model Family:          Poisson  Df Model:                  3
Link Function:          log      Scale:                  1.0000
Method:                 IRLS     Log-Likelihood:         -145.96
Date:                  Mon, 12 Apr 2021    Deviance:              194.06
Time:                  13:30:54    Pearson chi2:          178.
No. Iterations:         6
Covariance Type:        nonrobust
=====

```

	coef	std err	z	P> z	[0.025	0.975]
type	-0.2237	0.048	-4.693	0.000	-0.317	-0.130
construction	0.3714	0.060	6.231	0.000	0.255	0.488
operation	0.7680	0.103	7.471	0.000	0.567	0.969
months	8.095e-05	2.84e-06	28.487	0.000	7.54e-05	8.65e-05

```

=====

```

● 3.

Python Kernel

Python Kernel (Please execute the R-Kernel model below:

```
In [48]: # Package implementation
import pandas as pd
import numpy as np
import statsmodels.api as sm
from statsmodels.graphics.gofplots import ProbPlot
import matplotlib.pyplot as plt
import pylab
import seaborn as sns
```

```
In [49]: df = pd.read_csv("softdrink.csv")
df.insert(df.shape[1], 'Interept', 1)
```

```
In [50]: Y = df["Time"]
X = df.drop(['Time'], axis=1)
```

```
In [51]: multi_lin_reg = sm.OLS(Y, X).fit()
print(multi_lin_reg.summary())
```

```
In [52]: #residual standard error/residual standard deviation
a = abs(Y-multi_lin_reg.predict(X)).values
a = a**2
(sum(a)/22)**0.5
```

```
Out[52]: 3.2594734475800964
```


```
In [53]: '''
a = abs(Y-multi_lin_reg.predict(X)).values
print(np.std(a, axis=None, dtype=None, out=None, ddof=3, keepdims=np._NoValue))
print((sum((a-np.mean(a))*2)/(np.shape(a)[0]-3))*(0.5))
'''
```

```
Out[53]: '\na = abs(Y-multi_lin_reg.predict(X)).values\nprint(np.std(a, axis=None, dtype=
((a-np.mean(a))*2)/(np.shape(a)[0]-3))*(0.5))\n'
```

```
In [54]: # pvalues for parameters
multi_lin_reg.pvalues
```

```
Out[54]: Cases      3.254932e-09
Distance    6.312469e-04
Interept    4.417012e-02
dtype: float64
```

```
In [55]: sm.qqplot(Y-multi_lin_reg.predict(X).values, line='q')
pylab.show()
```




```
In [56]: # model values
model_fitted_y = multi_lin_reg.fittedvalues
# model residuals
model_residuals = multi_lin_reg.resid
# normalized residuals
model_norm_residuals = multi_lin_reg.get_influence().resid_studentized_internal
# absolute squared normalized residuals
model_norm_residuals_abs_sqrt = np.sqrt(np.abs(model_norm_residuals))
# absolute residuals
model_abs_resid = np.abs(model_residuals)
# leverage, from statsmodels internals
model_leverage = multi_lin_reg.get_influence().hat_matrix_diag
# cook's distance, from statsmodels internals
model_cooks = multi_lin_reg.get_influence().cooks_distance[0]
```

```
In [57]: plot_lm_1 = plt.figure()
plot_lm_1.axes[0] = sns.residplot(model_fitted_y, Y, data=df, lowess=True,
                                scatter_kws={'alpha': 0.5}, line_kws={'color': 'red', 'lw': 1, 'alpha': 0.8})
plot_lm_1.axes[0].set_title('Residuals vs Fitted')
plot_lm_1.axes[0].set_xlabel('Fitted values')
plot_lm_1.axes[0].set_ylabel('Residuals');
```

c:\users\andy\appdata\local\programs\python\python37\lib\site-packages\seaborn_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning




```
In [59]: plot_lm_4 = plt.figure()
plt.scatter(model_leverage, model_norm_residuals, alpha=0.5)
sns.regplot(model_leverage, model_norm_residuals, scatter=False, ci=False, lowess=True, line_kws={'color': 'red', 'lw': 1, 'alpha': 0.8});
plot_lm_4.axes[0].set_xlim(0, max(model_leverage)+0.01)
plot_lm_4.axes[0].set_ylim(-3, 5)
plot_lm_4.axes[0].set_title('Residuals vs Leverage')
plot_lm_4.axes[0].set_xlabel('Leverage')
plot_lm_4.axes[0].set_ylabel('Standardized Residuals');
```

```
# annotations
leverage_top_3 = np.flip(np.argsort(model_cooks), 0)[:3]
for i in leverage_top_3:
    plot_lm_4.axes[0].annotate(i+1, xy=(model_leverage[i], model_norm_residuals[i]));
```

c:\users\andy\appdata\local\programs\python\python37\lib\site-packages\seaborn_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning



R Kernel

R Kernel

```
In [1]: # read data
data = read.csv("softdrink.csv")
data
```

A data frame: 25 x 3

Time	Cases	Distance
<dbl>	<int>	<int>
16.68	7	560
11.50	3	220
12.03	3	340

```
In [2]: fit = lm(Time ~ Cases + Distance, data=data)
summary(fit)
```

```
Call:
lm(formula = Time ~ Cases + Distance, data = data)


Residuals:
    Min       1Q   Median       3Q      Max
-5.7880 -0.6629  0.4364  1.1566  7.4197

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  2.341231   1.096730   2.135  0.044170 *
Cases        1.615907   0.170735   9.464 3.25e-09 ***
Distance     0.014385   0.003613   3.981 0.000631 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Residual standard error: 3.259 on 22 degrees of freedom
Multiple R-squared:  0.9596,    Adjusted R-squared:  0.9559
F-statistic: 261.2 on 2 and 22 DF,  p-value: 4.687e-16
```

$$\text{Residual standard error} = \sqrt{\frac{\sum_{i=1}^n (Y_i - \hat{Y}_i)^2}{df}}$$

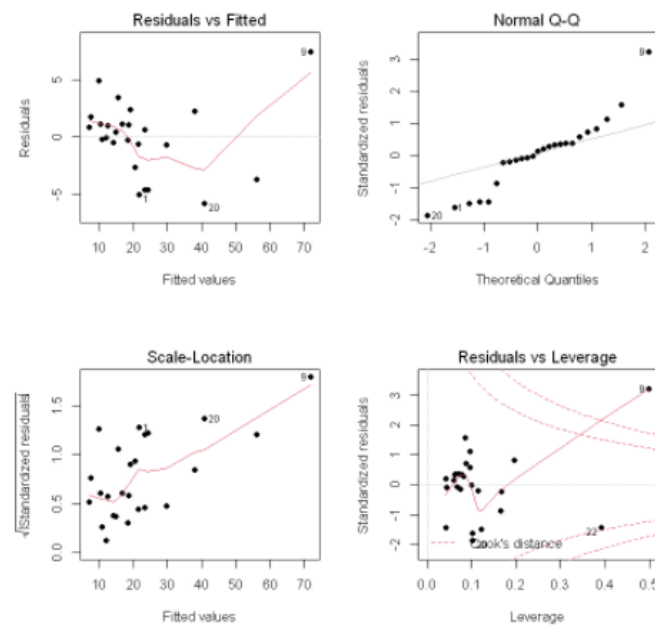
```
In [6]: res = resid(fit)
plot(fitted(fit), res)
abline(0,0)
```



```
In [3]: hist(residuals(fit))
```



```
In [4]: par(mfrow=c(2,2))
plot(fit,pch=16)
```



```
In [5]: influence.measures(fit)
```

```
Influence measures of
lm(formula = Time ~ Cases + Distance, data = data) :
```

	dfb.1_	dfb.Cass	dfb.Dstn	dffit	cov.r	cook.d	hat	inf
1	-0.18727	0.41131	-0.43486	-0.5709	0.871	1.00e-01	0.1018	
2	0.08979	-0.04776	0.01441	0.0986	1.215	3.38e-03	0.0707	
3	-0.00352	0.00395	-0.00285	-0.0052	1.276	9.46e-06	0.0987	
4	0.45196	0.08828	-0.27337	0.5008	0.876	7.76e-02	0.0854	
5	-0.03167	-0.01330	0.02424	-0.0395	1.240	5.43e-04	0.0750	
6	-0.01468	0.00179	0.00108	-0.0188	1.200	1.23e-04	0.0429	
7	0.07807	-0.02228	-0.01102	0.0790	1.240	2.17e-03	0.0818	
8	0.07120	0.03338	-0.05382	0.0938	1.206	3.05e-03	0.0637	
9	-2.57574	0.92874	1.50755	4.2961	0.342	3.42e+00	0.4983	*
10	0.10792	-0.33816	0.34133	0.3987	1.305	5.38e-02	0.1963	
11	-0.03427	0.09253	-0.00269	0.2180	1.172	1.62e-02	0.0861	
12	-0.03027	-0.04867	0.05397	-0.0677	1.291	1.60e-03	0.1137	
13	0.07237	-0.03562	0.01134	0.0813	1.207	2.29e-03	0.0611	
14	0.04952	-0.06709	0.06182	0.0974	1.228	3.29e-03	0.0782	
15	0.02228	-0.00479	0.00684	0.0426	1.192	6.32e-04	0.0411	
16	-0.00269	0.06442	-0.08419	-0.0972	1.369	3.29e-03	0.1659	
17	0.02886	0.00649	-0.01570	0.0339	1.219	4.01e-04	0.0594	
18	0.24856	0.18973	-0.27243	0.3653	1.069	4.40e-02	0.0963	
19	0.17256	0.02357	-0.09897	0.1862	1.215	1.19e-02	0.0964	
20	0.16804	-0.21500	-0.09292	-0.6718	0.760	1.32e-01	0.1017	
21	-0.16193	-0.29718	0.33641	-0.3885	1.238	5.09e-02	0.1653	
22	0.39857	-1.02541	0.57314	-1.1950	1.398	4.51e-01	0.3916	*
23	-0.15985	0.03729	-0.05265	-0.3075	0.890	2.99e-02	0.0413	
24	-0.11972	0.40462	-0.46545	-0.5711	0.948	1.02e-01	0.1206	
25	-0.01682	0.00085	0.00559	-0.0176	1.231	1.08e-04	0.0666	

● 4.

None

● 5.

```

In [1]: #Import libraries
        from scipy.stats import expon
        import math
        import numpy as np
        import matplotlib.pyplot as plt

In [2]: #Set the seed
        np.random.seed(44)

In [3]: # initial values
        x = 1

In [4]: # simulation parameters
        N = 1000
        burnin = 300

In [5]: #Using a list store the x_i and y_i
        X_Sample = []
        Y_Sample = []

In [6]: for i in range(N+burnin):
        #Sampling y for given x
        y = float(np.random.exponential(1+x, 1))
        x = float(np.random.exponential(1+y, 1))

        #Upgrading
        if i >= burnin:
            #Store the data
            X_Sample.append(x)
            Y_Sample.append(y)


```

```

In [7]: synth_plot = plt.scatter(X_Sample, Y_Sample,s=2)
        plt.xlabel("x")
        plt.ylabel("y")

```

Out[7]: Text(0, 0.5, 'y')



```

In [8]: sum = 0
        for i in range(N):
            x = X_Sample[i]
            y = Y_Sample[i]
            k = x+y+x*y
            try:
                s = 1/(math.exp(k))
            except OverflowError:
                s = 0
            sum = sum + s
        print(sum/1000)

```

0.06896669414414609


```

In [9]: def Gibbs_Sampling(N, burnin, initial_value, seed=None):
        if seed!=None:
            np.random.seed(seed)

        X_Sample = []
        Y_Sample = []
        x = initial_value
        for i in range(N+burnin):
            #Sampling y for given x
            y = float(np.random.exponential(1+x, 1))
            x = float(np.random.exponential(1+y, 1))

            #Upgrading
            if i >= burnin:
                #Store the data
                X_Sample.append(x)
                Y_Sample.append(y)

        sum = 0
        for i in range(N):
            x = X_Sample[i]
            y = Y_Sample[i]
            k = x*y+x*y
            try:
                s = 1/(math.exp(k))
            except OverflowError:
                s = 0
            sum = sum + s
        return sum/N


```

```

In [10]: C = []
        N = 1000
        burnin = 100
        initial_value = 10
        for i in range(N):
            C.append(Gibbs_Sampling(N, burnin, initial_value))
        synth_plot = plt.scatter([i for i in range(N)], C, s=1)
        plt.xlabel("the i-th sample")
        plt.ylabel("the value of c^-1")
        print("Average is: ", np.mean(C))

```

Average is: 0.06960987528553901



Thank you for reading my report.