

XinqianWang s4565489

### 1. (a)

My knowledge told me a regression MLP that uses identity activation functions for all neurons would have no difference with a OLS model. So I plan to build two Regression models using Pytorch and Sk-Learn respectively, and do iteration for 100 times to check if the things have happened.

```
import torch
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
import torch.optim as optim
import torch.nn as nn
from torch.nn.modules.loss import MSELoss
def regression_data(n=500, d=2):
    X = torch.rand(n, d)
    w = torch.rand(d+1)
    Y = X @ w[1:] + w[0] + torch.rand(n) * 0.1
    return X, Y
def ols3(X, Y, niter=200, lr=0.5):
    K = []
    Y = Y.reshape(-1, 1) #Change to n*1
    net = nn.Linear(2, 1)
    optimizer = optim.SGD(net.parameters(), lr=lr, momentum=0)
    mse = MSELoss()
    for i in range(niter):
        optimizer.zero_grad()
        loss = mse(net(X), Y)
        loss.backward()
        optimizer.step()
    for param in net.parameters():
        K.append(param.data)
    return K
ols3(X, Y)
```

```

def Data_7703_1_a():
    X, Y = regression_data(d=2)

    reg = LinearRegression()
    reg.fit(X, Y)
    y_pred = reg.predict(X)
    SK_MSE = mean_squared_error(Y, y_pred)

    L = ols3(X, Y)
    coef = L[0]
    Intercept = L[1]
    tensor_y_predict = coef@X.T+0.2947
    coef = torch.tensor([[0.8489, 0.5036]], requires_grad=True)
    tensor_y_predict = tensor_y_predict[0]
    tensor_y_predict = tensor_y_predict.detach().numpy()
    NLP_MSE = mean_squared_error(Y, y_pred)

    if NLP_MSE - SK_MSE != 0:
        print("!!!")

for i in range(1, 101):
    Data_7703_1_a()

```

I failed to output a better value than OLS, my understanding is that normally, it would not output lower MSE, may be in some extreme situation resulted by computer's calculation it would have a little difference. Basically the data scientist's announcement is not right.

1. (b)

$$\begin{aligned}
 \omega_{t+1} &= \omega_t - \eta_t \nabla L_{\lambda}(\omega_t) \\
 \nabla L_{\lambda}(\omega_t) &= \nabla L(\omega_t) + \lambda \omega_t \\
 \nabla L_{\lambda}(\omega_t) &> \nabla L(\omega_t) \\
 \omega_{t+1} &= \omega_t - \eta_t (\nabla L(\omega_t) + \lambda \omega_t) = (1 - \lambda) \omega_t - \eta_t \nabla L(\omega_t)
 \end{aligned}$$

We can see by using the L2 regularization the new weight  $t+1$  would decrease, which is equivalent to first multiply  $\omega$  by a constant value  $\lambda$ .

Compared to  $L(\omega)$ , the new loss function would decrease more in a gradient descent step. And we can also observe the weight decrease by  $\lambda$  times.

1. (c)

First: Suppose all elements are identical  $o_1 = o_2 = o_3 = \dots = o_c$   
 $\text{softmax}(o_1, \dots, o_c) = (\frac{1}{c}, \frac{1}{c}, \dots, \frac{1}{c}) = \text{softmax}_\beta(o_1, \dots, o_c)$

Second: We increase <sup>the value of</sup>  $o_i \in \{o_1, o_2, \dots, o_c\}$  to  $ko_i$  ( $k > 1$ ) and leave others do not change.  $\text{softmax}_\beta o_i = \frac{e^{\beta ko_i}}{e^{\beta o_1} + e^{\beta o_2} + \dots + e^{\beta ko_i} + \dots + e^{\beta o_c}}$   $\beta > 0$

When  $\beta$  increases the value of  $\text{softmax}_\beta o_i$  increases  
 Because  $\text{softmax}_\beta o_i = \frac{e^{\beta o_i} + e^{(k-1)\beta o_i}}{(c-1)e^{\beta o_i} + e^{(k-1)\beta o_i}} = \frac{e^{(k-1)\beta o_i}}{(c-1) + e^{(k-1)\beta o_i}}$

Let  $e^{(k-1)\beta o_i} = t \Rightarrow$  we get  $\frac{t}{c-1+t} = \frac{1+(t-1)}{c+(t-1)}$   
 $\begin{cases} k-1 > 0 \\ \beta > 0 \\ o_i \in \mathbb{R}^c \end{cases}$  so,  $t > 1$ ,  $\therefore$  as  $\beta$  increases  $\frac{t}{c-1+t}$  increases.

1. (d)

The  $w_1, w_2$  must meet 4 criterions below:

- ①  $(x_1, x_2) = (0, 0)$   $0 \times w_1 + 0 \times w_2 + b \leq 0$   
 $b \leq 0$
- ②  $(x_1, x_2) = (0, 1)$   $0 \times w_1 + 1 \times w_2 + b > 0 \rightarrow w_2 + b > 0 \rightarrow w_2 > 0$
- ③  $(x_1, x_2) = (1, 0)$   $1 \times w_1 + 0 \times w_2 + b > 0 \rightarrow w_1 + b > 0 \rightarrow w_1 > 0$
- ④  $(x_1, x_2) = (1, 1)$   $1 \times w_1 + 1 \times w_2 + b \leq 0 \rightarrow w_1 + w_2 + b \leq 0$

Set  $H_1$  for all possible from ①  
 $H_2$  for  $\dots$  from ②  
 $H_3$   $\dots$  ③  
 $H_4$   $\dots$  ④

$\{(w_1, w_2, b) \mid \forall w_1 \in H_3 \cap H_4, \forall w_2 \in H_2 \cap H_4, \forall b \in H_1 \cap H_2 \cap H_3 \cap H_4\} = \emptyset$

## 2. (b),(c)

```
In [ ]: def predict(self, X):
        return np.argmax(self.predict_proba(X), axis=1)

def predict_proba(self, X):
    X = np.dot(X, self.coef_.T) + self.intercept_
    e_X = np.exp(X - np.max(X))
    return softmax(e_X, axis=1)
```

```
def fit(self, X, y, lr=0.01, momentum=0, niter=1000):
    scaler = StandardScaler()
    X = scaler.fit_transform(X)
    self.classes_ = np.unique(y)
    self.class2int = dict((c, i) for i, c in enumerate(self.classes_))
    y = np.array([self.class2int[c] for c in y])

    n_features = X.shape[1]
    n_classes = len(self.classes_)

    self.intercept_ = np.zeros(n_classes)
    self.coef_ = np.zeros((n_classes, n_features))

    # Implement your gradient descent training code here: uncomment the code below to do "random training"
    self.intercept_ = np.random.randn(*self.intercept_.shape)
    self.coef_ = np.random.randn(*self.coef_.shape)

    for n in range(0, niter):
        #for n in range(0, 1):

        P_proba = self.predict_proba(X)
        y_predict = self.predict(X)

        for i in range(0, len(y)):
            if y[i] != y_predict[i]:
                rate = P_proba[i][y[i]]
                #print(y[i], y_predict[i])
                self.coef_ = self.coef_ - lr * (1/len(y)) * (1-rate) * X[i]
            else:
                ratej = P_proba[i][y_predict[i]]
                self.coef_ = self.coef_ - lr * (1/len(y)) * (0-ratej) * X[i]

        LOSS = log_loss(y, P_proba)
        acc_sum = 0.0
        acc_sum += (y_predict == y).sum()
        print(LOSS, "test accuracy: %f" % (acc_sum/len(y)))
    return self
```

My numpy function worked not well so I change to pytorch.

I deduced because of I have not do Standardlization to X, the distribution of X would got a lot of same value at first, but the arg max seems not functioned well on them. I think this is the reason. ¶

## 2. (d),(e),(f),(g)

```
1 [76]: import numpy as np
import torch.nn as nn
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.datasets import fetch_covtype
from sklearn import linear_model
from sklearn.preprocessing import StandardScaler

import torch
import torch.optim as optim
import torch.utils.data as Data

class LogisticRegression(nn.Module):
    def __init__(self, X,y):
        super(LogisticRegression, self).__init__()
        scaler = StandardScaler()
        X = scaler.fit_transform(X)
        K_train = np.max(X)
        X = X - K_train

        self.classes_ = np.unique(y)
        self.class2int = dict((c, i) for i, c in enumerate(self.classes_))
        y = np.array([self.class2int[c] for c in y])
        print(np.unique(y))
        n_features = X.shape[1]
        n_classes = len(self.classes_)

        self.w = nn.Parameter(torch.randn(n_classes, n_features))
        self.b = nn.Parameter(torch.zeros(n_classes))

        self._X = torch.from_numpy(X).type(torch.FloatTensor)
        self._y = torch.from_numpy(y).type(torch.LongTensor)

        self.net = nn.Sequential(
            nn.LogSoftmax()
        )
```

```

def fit(self, lr=10, momentum=0.9, niter=100, BATCH_SIZE=100):
    LOSS_FUNC = nn.CrossEntropyLoss()
    OPTIMIZER = torch.optim.SGD([self.w, self.b], lr=lr, momentum=momentum)
    train_set = Data.TensorDataset(self._X, self._y)
    train_loader = Data.DataLoader(dataset=train_set, batch_size=BATCH_SIZE, shuffle=True)
    for epoch in range(1, niter+1):
        loss_sum = 0.0
        for step, (x, y) in enumerate(train_loader):
            y_pred = self.predict_proba(x)
            y = y.squeeze()
            loss = LOSS_FUNC(y_pred, y)
            loss_sum += loss
            OPTIMIZER.zero_grad()
            loss.backward()
            OPTIMIZER.step()
        print("epoch: %d, loss: %f" % (epoch, loss_sum/BATCH_SIZE))

def predict_proba(self, X):
    X = torch.mm(X, self.w.T) + self.b.T
    return self.net(X)

def predict(self, X):
    X = self.predict_proba(X)
    return X.argmax(dim=1)

if __name__ == '__main__':
    X, y = fetch_covtype(return_X_y=True)
    X_tr, X_ts, y_tr, y_ts = train_test_split(X, y, test_size=0.3, random_state=5)

    LR = LogisticRegression(X_tr, y_tr)
    LR.fit()

```

epoch: 1, loss: 4526806016.000000  
epoch: 2, loss: 4724474880.000000  
epoch: 3, loss: 4551279616.000000  
epoch: 4, loss: 4540598272.000000  
epoch: 5, loss: 4492083712.000000  
epoch: 6, loss: 4432264704.000000  
epoch: 7, loss: 4452489728.000000  
epoch: 8, loss: 4591495680.000000  
epoch: 9, loss: 4366866432.000000  
epoch: 10, loss: 4382897152.000000  
epoch: 11, loss: 4477735936.000000  
epoch: 12, loss: 4415218688.000000  
epoch: 13, loss: 4277500928.000000  
epoch: 14, loss: 4345388544.000000  
epoch: 15, loss: 4265158912.000000  
epoch: 16, loss: 4466298880.000000  
epoch: 17, loss: 4231500800.000000  
epoch: 18, loss: 4175854848.000000  
epoch: 19, loss: 4022260224.000000  
epoch: 20, loss: 4108012544.000000  
epoch: 21, loss: 4263953664.000000  
epoch: 22, loss: 4055972608.000000  
epoch: 23, loss: 4102737152.000000  
epoch: 24, loss: 4008943360.000000  
epoch: 25, loss: 4011353088.000000  
epoch: 26, loss: 3923219200.000000  
epoch: 27, loss: 4133430784.000000  
epoch: 28, loss: 3895300352.000000  
epoch: 29, loss: 3743054336.000000  
epoch: 30, loss: 3813127936.000000  
epoch: 31, loss: 3868281088.000000  
epoch: 32, loss: 3760358400.000000  
epoch: 33, loss: 3684607488.000000  
epoch: 34, loss: 3874646272.000000  
epoch: 35, loss: 3870728192.000000  
epoch: 36, loss: 3851150592.000000  
epoch: 37, loss: 3693130240.000000  
epoch: 38, loss: 3664062464.000000  
epoch: 39, loss: 3643851520.000000  
epoch: 40, loss: 3683076352.000000  
epoch: 41, loss: 3533725696.000000  
epoch: 42, loss: 3451090944.000000  
epoch: 43, loss: 3455759360.000000

```
epoch: 44, loss: 3506829824.000000
epoch: 45, loss: 3553481728.000000
epoch: 46, loss: 3481617664.000000
epoch: 47, loss: 3395985664.000000
epoch: 48, loss: 3374000640.000000
epoch: 49, loss: 3395615232.000000
epoch: 50, loss: 3351093248.000000
epoch: 51, loss: 3345948416.000000
epoch: 52, loss: 3329639168.000000
epoch: 53, loss: 3258699520.000000
epoch: 54, loss: 3150112768.000000
epoch: 55, loss: 3108879104.000000
epoch: 56, loss: 3281637888.000000
epoch: 57, loss: 3100217600.000000
epoch: 58, loss: 3155966720.000000
epoch: 59, loss: 3046773760.000000
epoch: 60, loss: 3140528384.000000
epoch: 61, loss: 3079246848.000000
epoch: 62, loss: 3042024704.000000
epoch: 63, loss: 3019451392.000000
epoch: 64, loss: 2949055488.000000
epoch: 65, loss: 2839685376.000000
epoch: 66, loss: 2882721536.000000
epoch: 67, loss: 2838137856.000000
epoch: 68, loss: 2980979712.000000
epoch: 69, loss: 2876399872.000000
epoch: 70, loss: 2796915200.000000
epoch: 71, loss: 2720542464.000000
epoch: 72, loss: 2791089664.000000
epoch: 73, loss: 2588010752.000000
epoch: 74, loss: 2650686464.000000
epoch: 75, loss: 2723838720.000000
epoch: 76, loss: 2809645312.000000
epoch: 77, loss: 2575085056.000000
epoch: 78, loss: 2678593536.000000
epoch: 79, loss: 2402688512.000000
epoch: 80, loss: 2682606592.000000
epoch: 81, loss: 2428070912.000000
epoch: 82, loss: 2644584960.000000
epoch: 83, loss: 2479998720.000000
epoch: 84, loss: 2373795072.000000
epoch: 85, loss: 2491685376.000000
epoch: 86, loss: 2525163776.000000
epoch: 87, loss: 2361658368.000000
epoch: 88, loss: 2352394496.000000
```



```

: y_ts = y_ts-1
scaler = StandardScaler()
X_ts = scaler.fit_transform(X_ts)
X_ts = torch.from_numpy(X_ts).type(torch.FloatTensor)
y_ts = torch.from_numpy(y_ts).type(torch.LongTensor)
acc_sum = 0.0
acc_sum += (LR.predict(X_ts) == y_ts.squeeze()).sum()
print("test accuracy: %f" %(acc_sum/len(y_ts)))

```

test accuracy: 0.594811

C:\Users\Andy\Anaconda3\lib\site-packages\torch\nn\modules\container.py:117: UserWarning: Implicit dimension choice for log\_softmax has been deprecated. Change the call to include dim=X as an argument.  
input = module(input)

### 3. (b),(c)

Sorry for late submission, I unconsciously delete the code for (b) and (c), I just need to redo again. Blew is the redo answer, for me, may be a little messy.

```

2]: import numpy as np

from sklearn.base import clone
from sklearn.datasets import load_boston
from sklearn.metrics import mean_squared_error
from sklearn.linear_model import LinearRegression, RANSACRegressor, TheilSenRegressor
from sklearn.model_selection import train_test_split
from sklearn.utils import check_random_state

def corrupt(X, y, outlier_ratio=0.1, random_state=None):
    random = check_random_state(random_state)

    n_samples = len(y)
    n_outliers = int(outlier_ratio*n_samples)

    W = X.copy()
    z = y.copy()

    mask = np.ones(n_samples).astype(bool)
    outlier_ids = random.choice(n_samples, n_outliers)
    mask[outlier_ids] = False

    W[mask, 4] *= 0.1

    return W, z

```

```

class ENOLS:
    def __init__(self, n_estimators=100, sample_size='auto'):
        """
        Parameters
        -----
        n_estimators: number of OLS models to train
        sample_size: size of random subset used to train the OLS models, default to 'auto'
            - If 'auto': use subsets of size n_features+1 during training
            - If int: use subsets of size sample_size during training
            - If float: use subsets of size ceil(n_sample*sample_size) during training
        """

        self.n_estimators = n_estimators
        self.sample_size = sample_size

```

```

def fit(self, X, y, random_state=None):
    """
    Train ENOLS on the given training set.

    Parameters
    -----
    X: an input array of shape (n_sample, n_features)
    y: an array of shape (n_sample,) containing the values for the input examples

    Return
    -----
    self: the fitted model
    """
    if self.sample_size == 'auto':
        S_size = X.shape[1]+1
    elif isinstance(self.sample_size, int):
        S_size = self.sample_size
    elif isinstance(self.sample_size, float):
        S_size = X.shape[0]*self.sample_size
    else:
        xxxxxxxx

    # use random instead of np.random to sample random numbers below
    random = check_random_state(random_state)

    # add all the trained OLS models to this list
    self.estimators_ = []

    # write your training code below. your code should support the
    # n_estimators and sample_size hyper-parameters described in the
    # documentation for the __init__ function
    self.base_estimator_ = LinearRegression()
    for n in range(0, self.n_estimators):
        estimator = clone(self.base_estimator_)
        X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=S_size, random_state=random)
        estimator.fit(X_train, y_train)
        self.estimators_.append(estimator)
    return self

```

```

def predict(self, X, method='average'):
    """
    Parameters
    -----
    X: an input array of shape (n_sample, n_features)
    method: 'median' or 'average', corresponding to predicting median and
           mean of the OLS models' predictions respectively.

    Returns
    -----
    y: an array of shape (n_samples,) containig the predicted values
    """

    if method == 'average':
        MEAN = []
        for SE in self.estimators_:
            P = SE.predict(X)
            MEAN.append(P)
        MEAN = np.array(MEAN)
        return np.mean(MEAN, axis=0)
    elif method == 'median':
        MEDIAN = []
        for SE in self.estimators_:
            P = SE.predict(X)
            MEDIAN.append(P)
        MEDIAN = np.array(MEDIAN)
        return np.median(MEDIAN, axis=0)
    else:
        xxxxxxxx

```

```

if __name__ == '__main__':
    X, y = load_boston(return_X_y=True)
    X_tr, X_ts, y_tr, y_ts = train_test_split(X, y, test_size=0.3, random_state=42)
    W, z = corrupt(X_tr, y_tr, outlier_ratio=0.1, random_state=42)
    '''

    reg = LinearRegression()
    reg.fit(X_tr, y_tr)
    print(mean_squared_error(y_ts, reg.predict(X_ts)))

    reg = LinearRegression()
    reg.fit(W, z)
    print(mean_squared_error(y_ts, reg.predict(X_ts)))

    EL = ENOLS()
    EL.fit(X_tr, y_tr)
    EL.predict(X_ts)
    '''

    OLD = []
    TSR = []
    ELS = []
    for i in range(0, 51):
        p = 0.01*i
        W, z = corrupt(X_tr, y_tr, outlier_ratio=p, random_state=42)

        LR = LinearRegression()
        TR = TheilSenRegressor()
        ES = ENOLS()

        LR.fit(W, z)
        TR.fit(W, z)
        ES.fit(W, z)


        OLD.append(mean_squared_error(y_ts, LR.predict(X_ts)))
        TSR.append(mean_squared_error(y_ts, TR.predict(X_ts)))
        ELS.append(mean_squared_error(y_ts, ES.predict(X_ts, method='median')))

```

### 3 (d)

**ENOLS seems to be less stable than OLD and TheilSenRegressor**


```
: import matplotlib.pyplot as plt
k = [i*0.01 for i in range(0,51)]
fig, ax = plt.subplots()
plt.plot(k, OLD, label="OLD")
plt.plot(k, TSR, label="TheilSenRegressor")
plt.plot(k, ELS, label="ENOLS")
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left', borderaxespad=0.)
ax.set_xlabel('number of k')
ax.set_ylabel('error')
plt.show()
```



```

: import matplotlib.pyplot as plt
k = [i*0.01 for i in range(0,51)]
fig, ax = plt.subplots()
plt.plot(k, OLD, label="OLD")
plt.plot(k, TSR, label="TheilSenRegressor")
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left', borderaxespad=0.)
ax.set_xlabel('number of k')
ax.set_ylabel('error')
plt.show()

```




3 (e)

**ENOLS still not stable than others and TheilSenRegressor Algorithm has the highest robustness.**

```

: import matplotlib.pyplot as plt
k = [i*0.01 for i in range(0,51)]
fig, ax = plt.subplots()
plt.plot(k, OLD, label="OLD")
plt.plot(k, TSR, label="TheilSenRegressor")
plt.plot(k, ELS, label="ENOLS")
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left', borderaxespad=0.)
ax.set_xlabel('number of k')
ax.set_ylabel('error')
plt.show()


```



### 3 (f)

When the number of estimators increases to a big number. The ENLOS gets to a highest stable condition compared to the other two algorithms.


```
[180]: import matplotlib.pyplot as plt
k = [i*0.01 for i in range(0,51)]
fig, ax = plt.subplots()
plt.plot(k, OLD, label="OLD")
plt.plot(k, TSR, label="TheilSenRegressor")
plt.plot(k, ELS, label="ENOLS")
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left', borderaxespad=0.)
ax.set_xlabel('number of k')
ax.set_ylabel('error')
plt.show()
```



### 3 (g)

The fluctuation of ENOLS tends to become a little more placid than before pictures, but the robustness seems to be not as good as the previous ENOLS model, probably because we decrease the number of subset.

```
: import matplotlib.pyplot as plt
k = [i*0.01 for i in range(0,51)]
fig, ax = plt.subplots()
plt.plot(k, OLD, label="OLD")
plt.plot(k, TSR, label="TheilSenRegressor")
plt.plot(k, ELS, label="ENOLS")
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left', borderaxespad=0.)
ax.set_xlabel('number of k')
ax.set_ylabel('error')
plt.show()
```



3(h)

$0 < q \leq 1$   $q$  is a constant

$m = nq$   $n$ , number of samples

$p$  rate of outliers

Assume we have  $n$  estimators, every do sampling, the rate of get an outliers is  $p$ . if we sample  $nq$  data for each set,

the possibility of getting an outliers is  ~~$(p)^m = p^{nq} = p^n$~~

$$(p)^m = p^{nq} \geq p^n$$

Because  $0 < p < 1$  and  $nq \leq n$

If  $n$  is large, for example  $1 \times 10^9$ ,  $p$  is 0.1, Then the possibility of we pick a outlier during sampling is very close to 0.

In other words sampling a fixed proportion  $q$  of the training set is just increasing the possibility of getting an outlier, which is really a bad idea.